How do actor-loss gradients compare?

Final 20k checkpoint · each point is one training seed · fixed retained logged probes

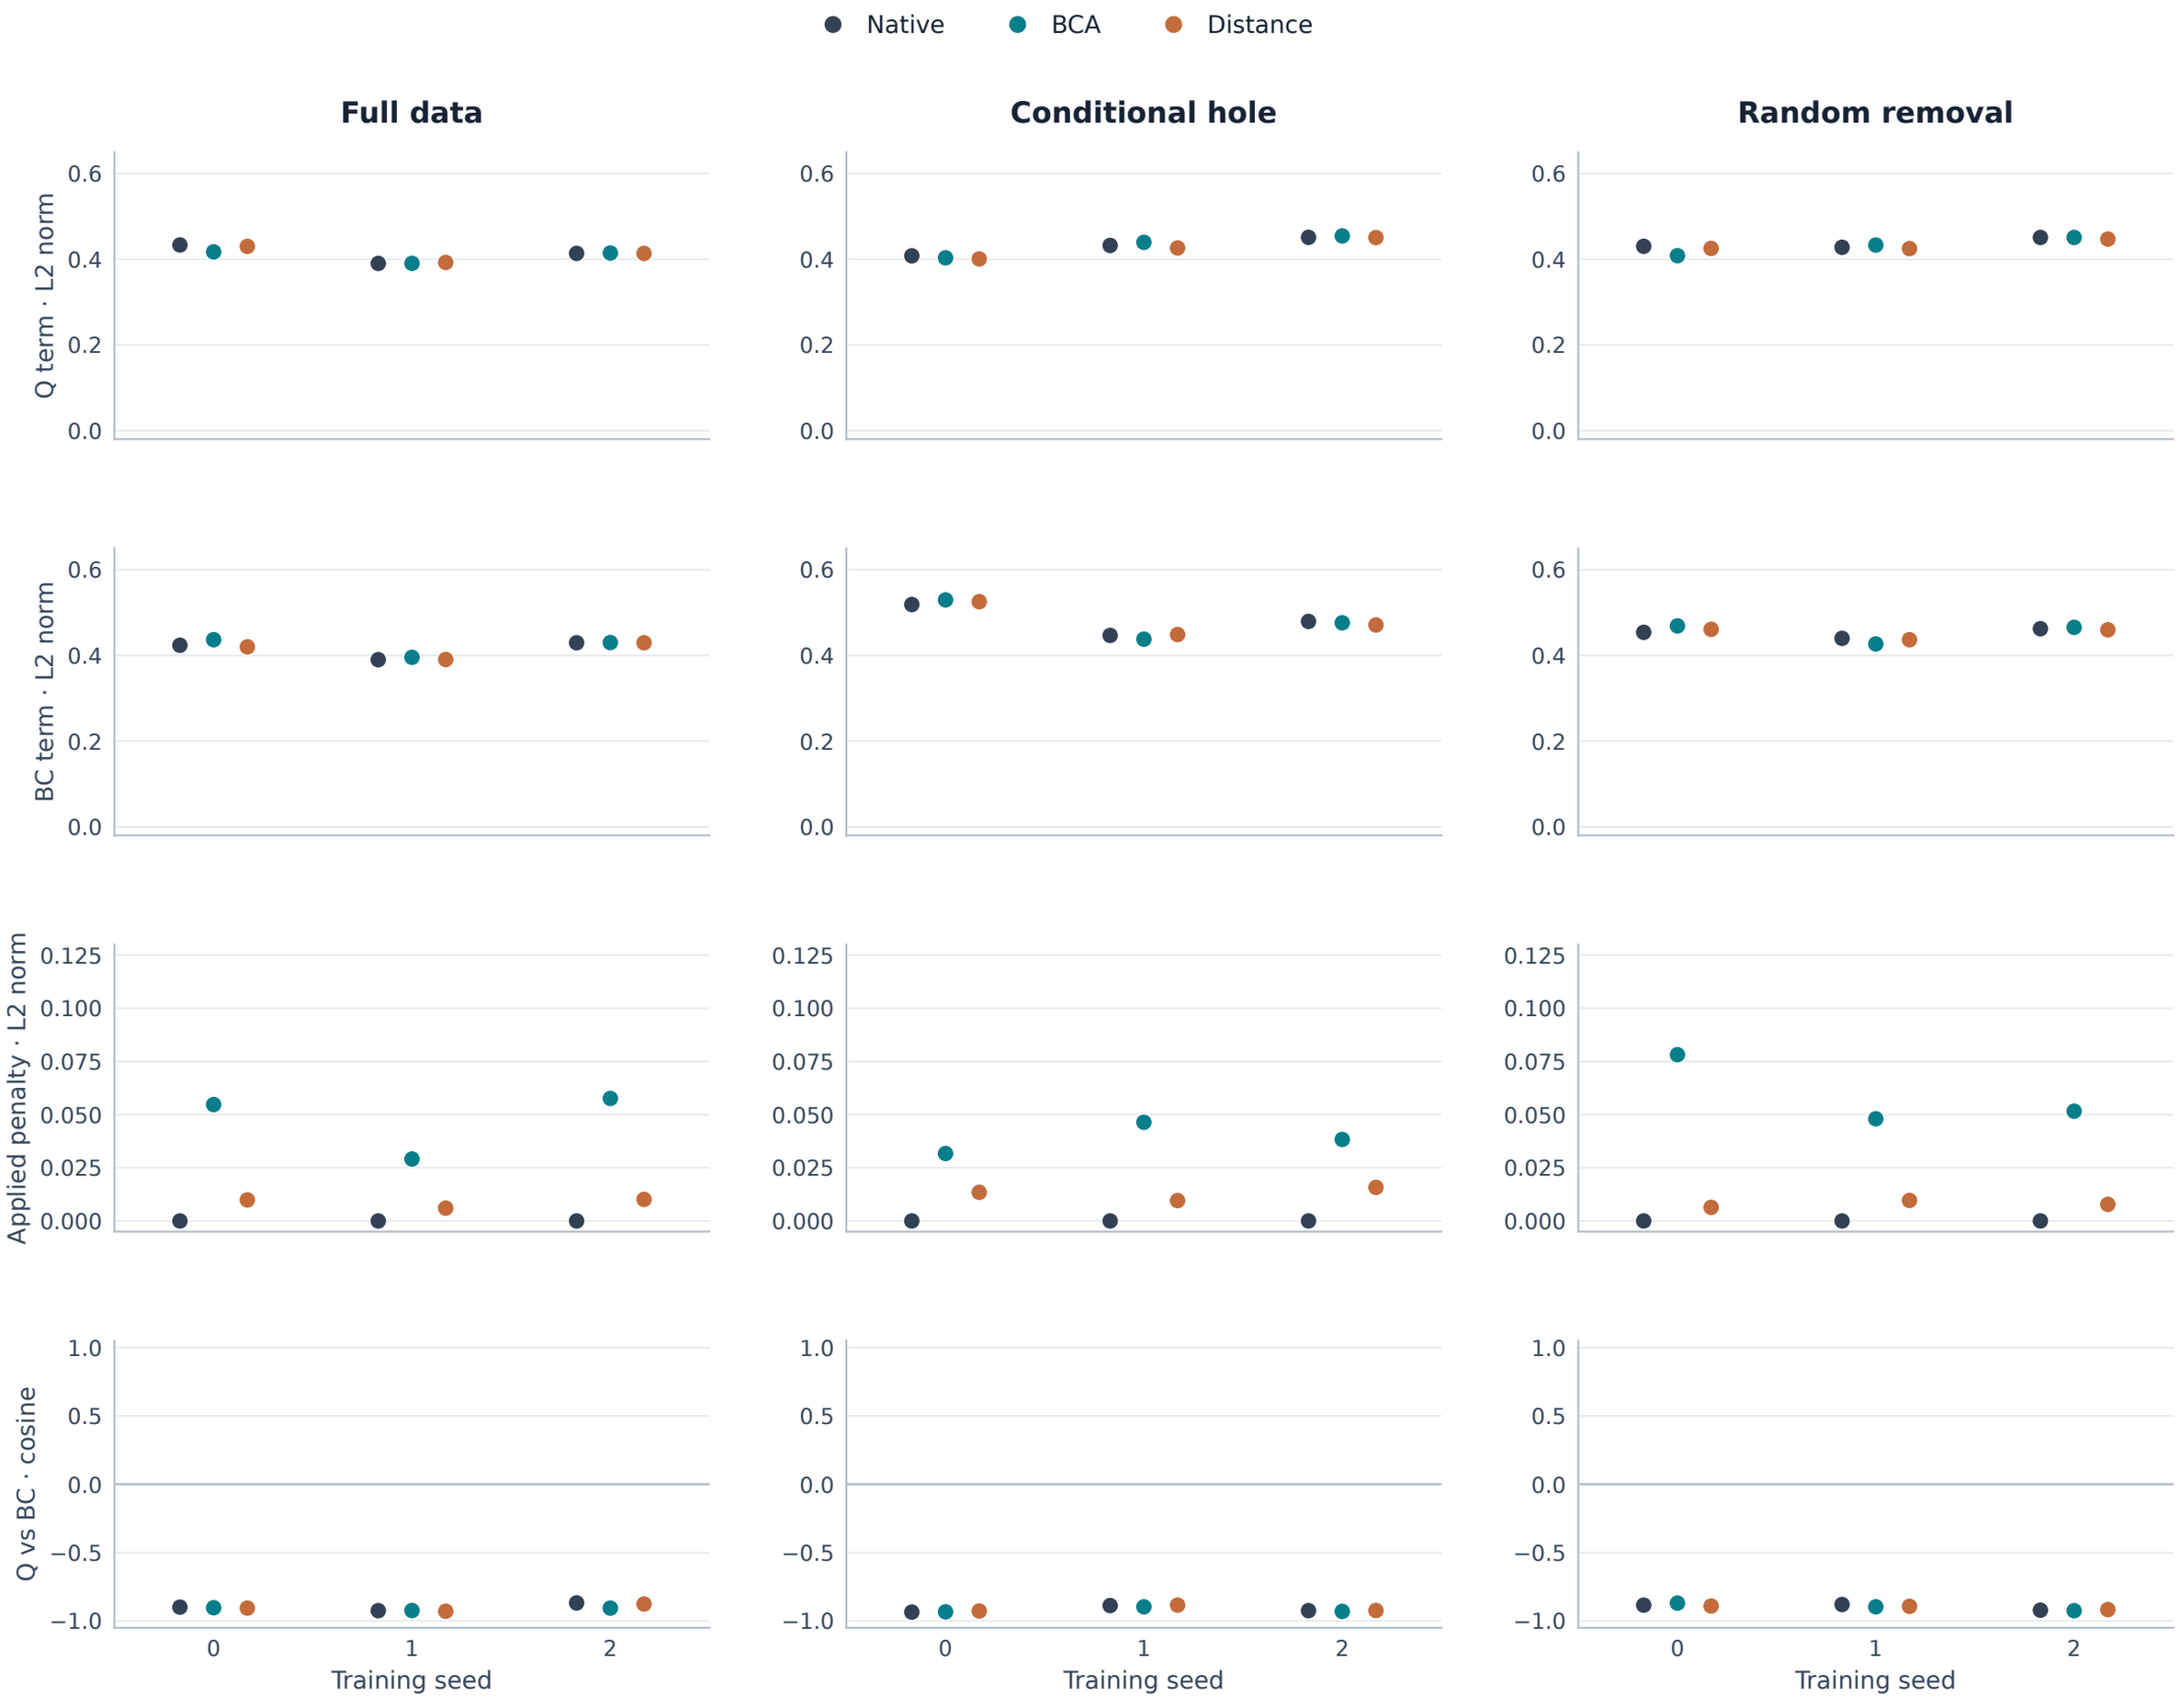

Norms and cosine are computed from each mean actor-parameter loss gradient across three fixed batches.
These are local gradients at one checkpoint, not executed updates or causal explanations of policy returns. Native has no applied penalty. Different panels use different y scales.