Does the BCA adjustment improve the same critic?

BCA-trained critics only | Same states, actions and own-policy reference for each pair | Final checkpoint

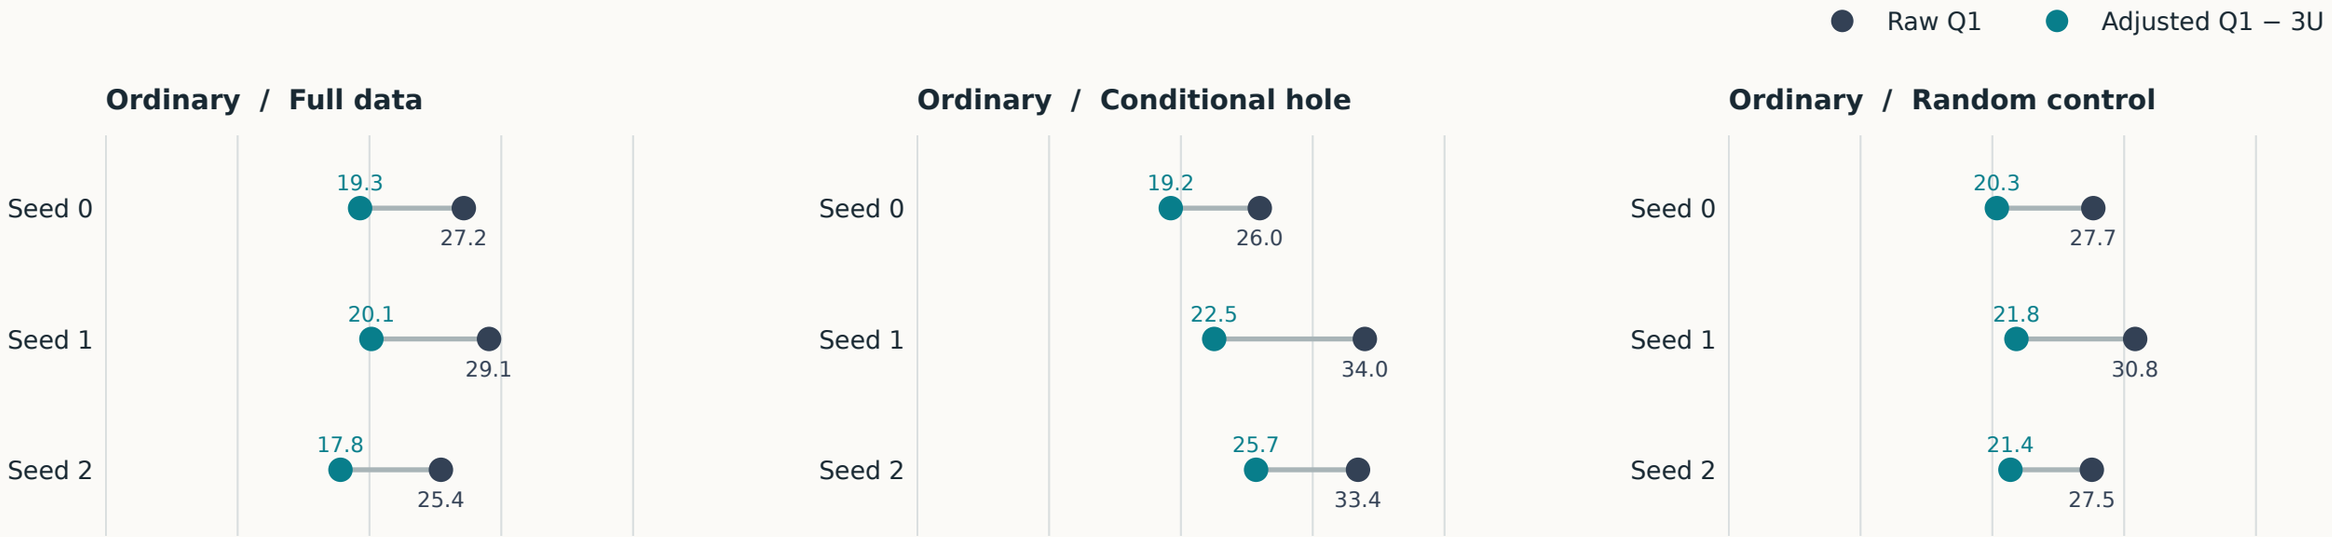

High-load fixed-band actions; smaller MAE is better. The strong-current, full-data seed 0 is a counterexample: 16.3 → 20.8.
21 ordinary / 32 strong-current anchor families per seed, with three correlated actions per anchor. This tests score adjustment, not policy-return causation.
MAE = mean absolute error in discounted-return units; U = BCA uncertainty score. Source: authenticated report 55cae667d369.