How do scores compare against one shared policy reference?

High-load fixed-band actions | Same states, actions and full-data native 20k reference per seed | Final checkpoint

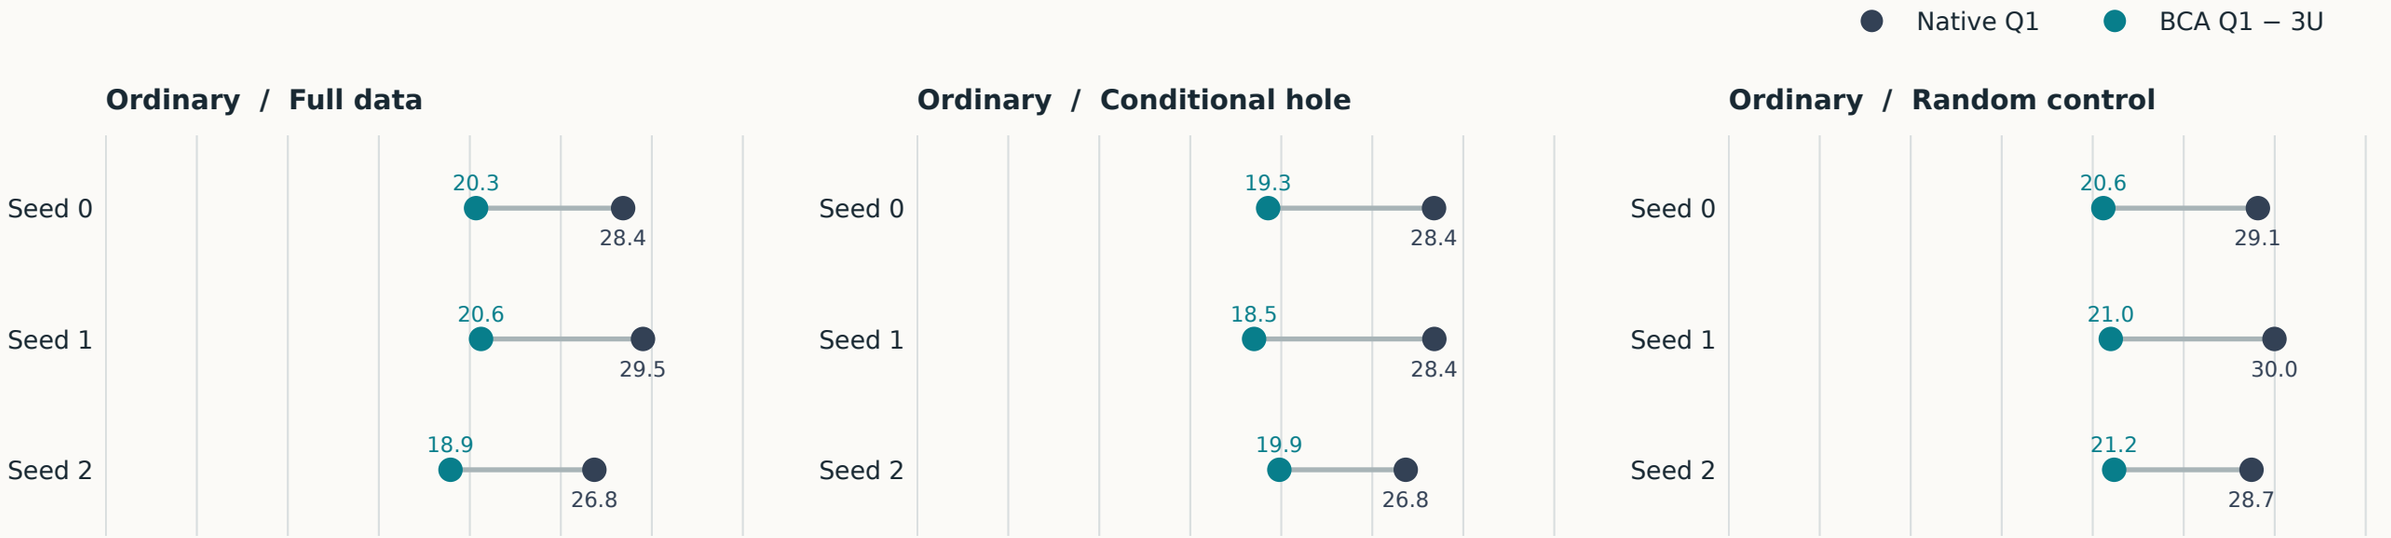

Reference is the full-data native 20k continuation for the same seed and context, not each model's own policy. This is deviation from a common reference, not operational accuracy for every policy. 21 ordinary / 32 strong-current anchors per seed; three correlated actions per anchor.
MAE = mean absolute error in discounted-return units; U = BCA uncertainty score. Source: authenticated report 55cae667d369.