Raw critic optimism and actor-score pessimism

Own-action probes: $Q_1 - Q\pi$ versus $Q_1 - \kappa U - Q\pi$. The actor score is not the learned critic prediction.

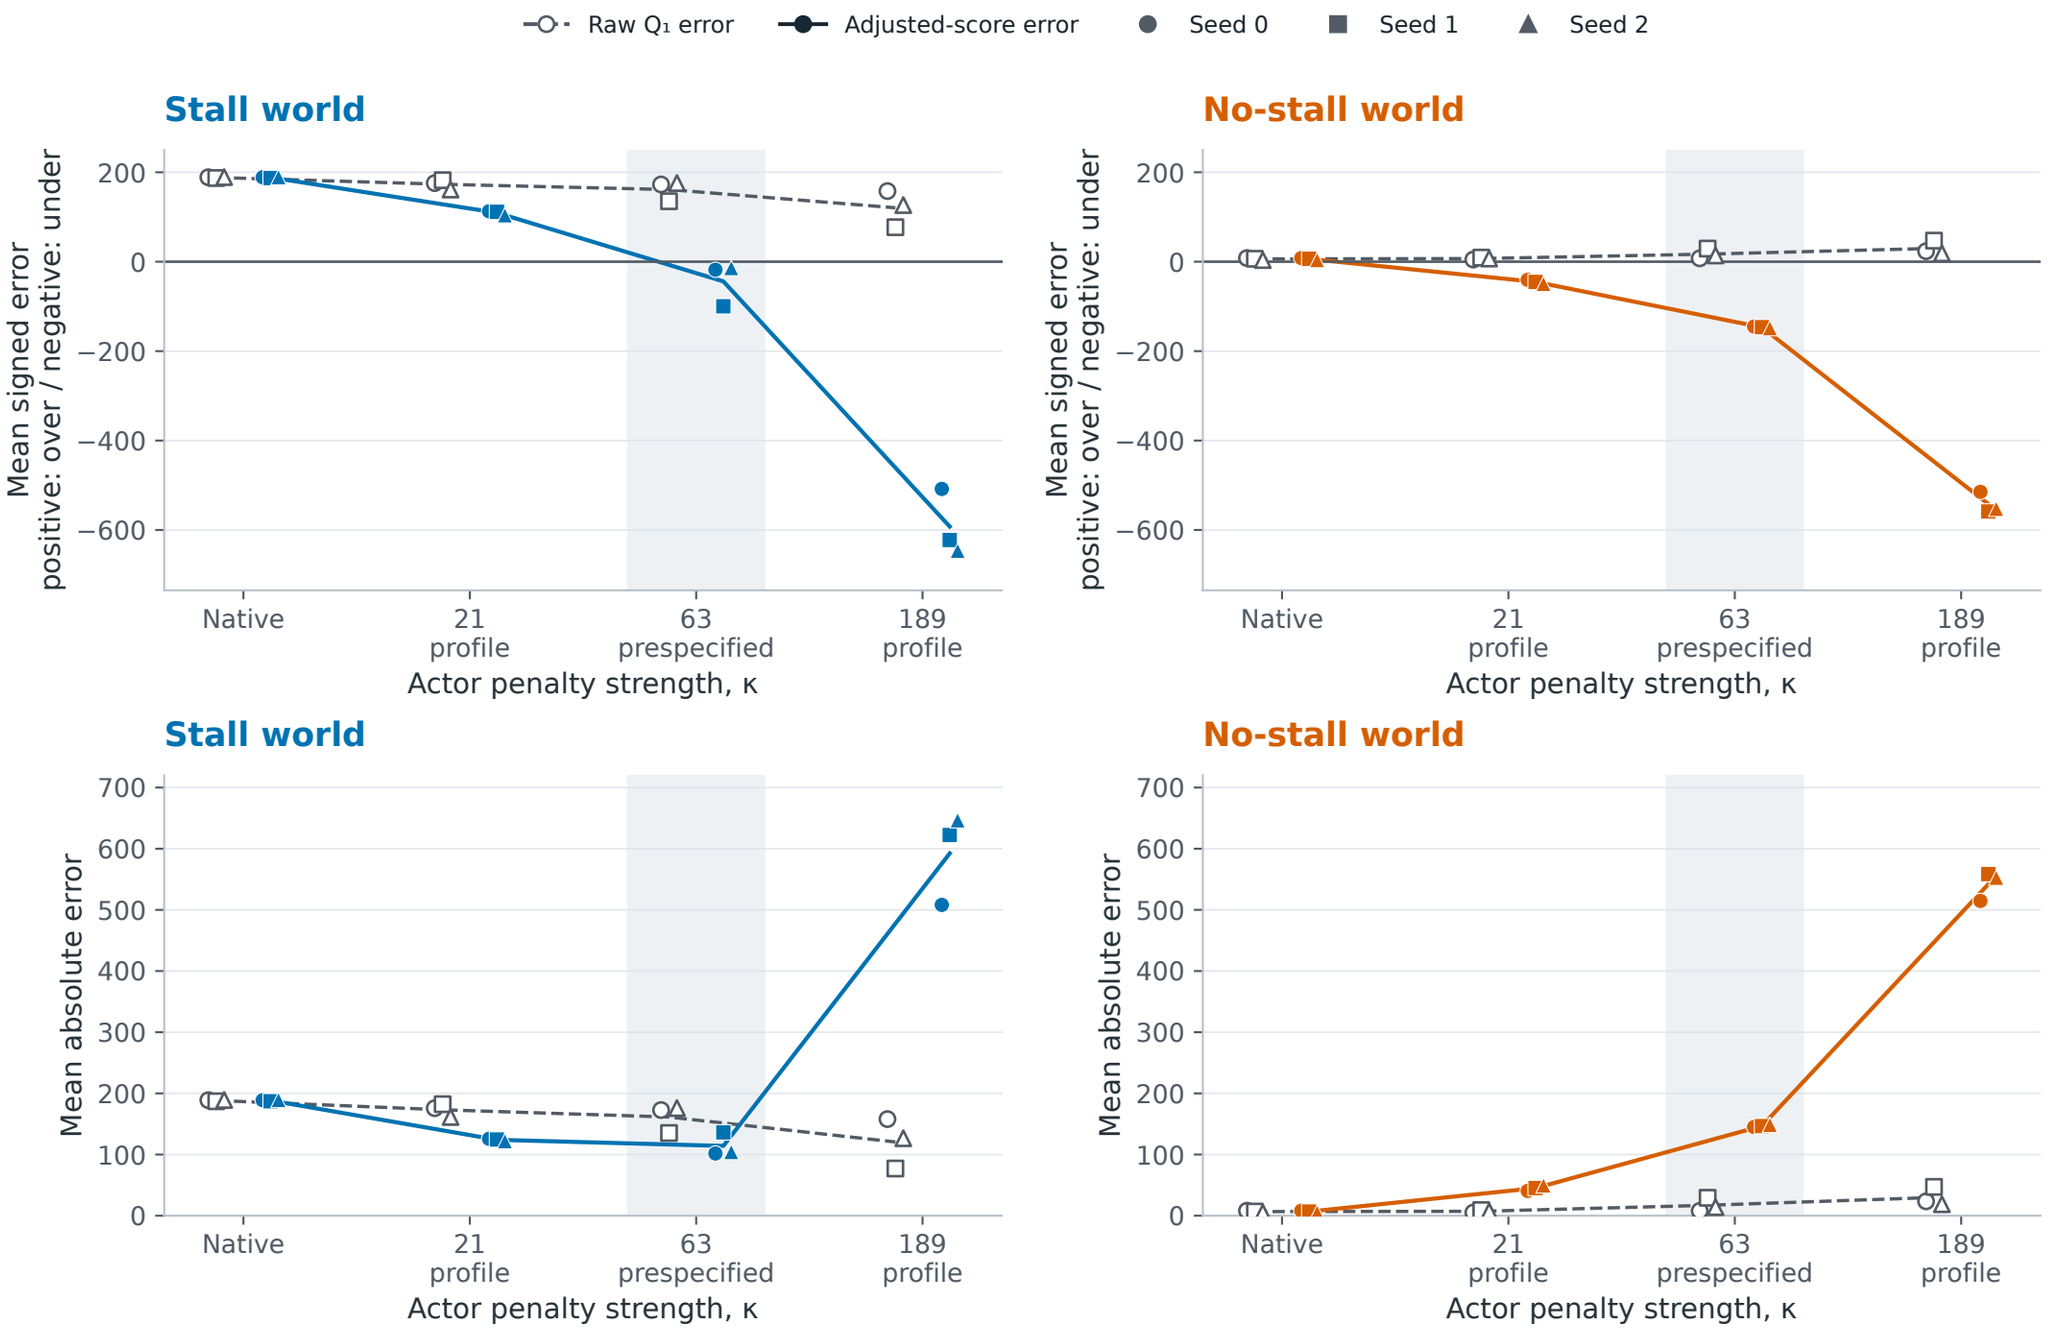

48 own-action probes per policy (16 at each of t = 0, 25, 50). Points: training seeds; lines: three-seed means.
Visited states and continuation policies differ across trained policies. These are not common-state critic-accuracy comparisons or true-Q coverage guarantees.